



RFP 132 – Asynchronous Services

Final

10 Pages

Abstract

The original requirements document for Distributed OSGi, RFP 88, included a requirement to support asynchronous communication protocols. This was deferred to the next release due to time constraints. However, this remains a significant requirement for enterprise applications using the OSGi Framework, especially to address application requirements for loose coupling, scalability, and reliability. This document expands on the requirements for asynchronous OSGi Services and how they would be added to the OSGi Framework. Asynchronous Services are expected to be a very elegant way to fit asynchronous distributed computing into the OSGi programming model.

Copyright © OSGi Alliance 2013.

All company, brand and product names contained within this document may be trademarks that are the sole property of the respective owners. The above notice must be included on all copies of this document that are made.

0 Document Information

0.1 License

DISTRIBUTION AND FEEDBACK LICENSE, Version 2.0

The OSGi Alliance hereby grants you a limited copyright license to copy and display this document (the "Distribution") in any medium without fee or royalty. This Distribution license is exclusively for the purpose of reviewing and providing feedback to the OSGi Alliance. You agree not to modify the Distribution in any way and further agree to not participate in any way in the making of derivative works thereof, other than as a necessary result of reviewing and providing feedback to the Distribution. You also agree to cause this notice, along with the accompanying consent, to be included on all copies (or portions thereof) of the Distribution. The OSGi Alliance also grants you a perpetual, non-exclusive, worldwide, fully paid-up, royalty free, limited license (without the right to sublicense) under any applicable copyrights, to create and/or distribute an implementation of the Distribution that: (i) fully implements the Distribution including all its required interfaces and functionality; (ii) does not modify, subset, superset or otherwise extend the OSGi Name Space, or include any public or protected packages, classes, Java interfaces, fields or methods within the OSGi Name Space other than those required and authorized by the Distribution. An implementation that does not satisfy limitations (i)-(ii) is not considered an implementation of the Distribution, does not receive the benefits of this license, and must not be described as an implementation of the Distribution. "OSGi Name Space" shall mean the public class or interface declarations whose names begin with "org.osgi" or any recognized successors or replacements thereof. The OSGi Alliance expressly reserves all rights not granted pursuant to these limited copyright licenses including termination of the license at will at any time.

EXCEPT FOR THE LIMITED COPYRIGHT LICENSES GRANTED ABOVE, THE OSGi ALLIANCE DOES NOT GRANT, EITHER EXPRESSLY OR IMPLIEDLY, A LICENSE TO ANY INTELLECTUAL PROPERTY IT, OR ANY THIRD PARTIES, OWN OR CONTROL. Title to the copyright in the Distribution will at all times remain with the OSGi Alliance. The example companies, organizations, products, domain names, e-mail addresses, logos, people, places, and events depicted therein are fictitious. No association with any real company, organization, product, domain name, email address, logo, person, place, or event is intended or should be inferred.

THE DISTRIBUTION IS PROVIDED "AS IS," AND THE OSGi ALLIANCE (INCLUDING ANY THIRD PARTIES THAT HAVE CONTRIBUTED TO THE DISTRIBUTION) MAKES NO REPRESENTATIONS OR WARRANTIES, EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, OR TITLE; THAT THE CONTENTS OF THE DISTRIBUTION ARE SUITABLE FOR ANY PURPOSE; NOR THAT THE IMPLEMENTATION OF SUCH CONTENTS WILL NOT INFRINGE ANY THIRD PARTY PATENTS, COPYRIGHTS, TRADEMARKS OR OTHER RIGHTS.

NEITHER THE OSGi ALLIANCE NOR ANY THIRD PARTY WILL BE LIABLE FOR ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF OR RELATING TO ANY USE OR DISTRIBUTION OF THE DISTRIBUTION.

Implementation of certain elements of this Distribution may be subject to third party intellectual property rights, including without limitation, patent rights (such a third party may or may not be a member of the OSGi Alliance). The OSGi Alliance is not responsible and shall not be held responsible in any manner for identifying or failing to identify any or all such third party intellectual property rights.

The Distribution is a draft. As a result, the final product may change substantially by the time of final publication, and you are cautioned against relying on the content of this Distribution. You are encouraged to update any implementation of the Distribution if and when such Distribution becomes a final specification.

The OSGi Alliance is willing to receive input, suggestions and other feedback ("Feedback") on the Distribution. By providing such Feedback to the OSGi Alliance, you grant to the OSGi Alliance and all its Members a non-exclusive, non-transferable,

worldwide, perpetual, irrevocable, royalty-free copyright license to copy, publish, license, modify, sublicense or otherwise distribute and exploit your Feedback for any purpose. Likewise, if incorporation of your Feedback would cause an implementation of the Distribution, including as it may be modified, amended, or published at any point in the future (“Future Specification”), to necessarily infringe a patent or patent application that you own or control, you hereby commit to grant to all implementers of such Distribution or Future Specification an irrevocable, worldwide, sublicenseable, royalty free license under such patent or patent application to make, have made, use, sell, offer for sale, import and export products or services that implement such Distribution or Future Specification. You warrant that (a) to the best of your knowledge you have the right to provide this Feedback, and if you are providing Feedback on behalf of a company, you have the rights to provide Feedback on behalf of your company; (b) the Feedback is not confidential to you and does not violate the copyright or trade secret interests of another; and (c) to the best of your knowledge, use of the Feedback would not cause an implementation of the Distribution or a Future Specification to necessarily infringe any third-party patent or patent application known to you. You also acknowledge that the OSGi Alliance is not required to incorporate your Feedback into any version of the Distribution or a Future Specification.

I HEREBY ACKNOWLEDGE AND AGREE TO THE TERMS AND CONDITIONS DELINEATED ABOVE.

0.2 Trademarks

OSGi™ is a trademark, registered trademark, or service mark of the OSGi Alliance in the US and other countries. Java is a trademark, registered trademark, or service mark of Oracle Corporation in the US and other countries. All other trademarks, registered trademarks, or service marks used in this document are the property of their respective owners and are hereby recognized.

0.3 Feedback

This document can be downloaded from the OSGi Alliance design repository at <https://github.com/osgi/design> The public can provide feedback about this document by opening a bug at <https://www.osgi.org/bugzilla/>.

0.4 Table of Contents

0 Document Information.....	2
0.1 License.....	2
0.2 Trademarks.....	3
0.3 Feedback.....	3
0.4 Table of Contents.....	3
0.5 Terminology and Document Conventions.....	4
0.6 Revision History.....	4
1 Introduction.....	4
2 Application Domain.....	6
Communicate using messages	8
2.1 Asynchronous programming models.....	8
2.2 Mixture of programming models.....	8
2.3 Terminology + Abbreviations.....	9
3 Problem Description.....	9
3.1 Asynchronous Services.....	9
4 Use Cases.....	9
5 Requirements.....	10
6 Document Support.....	11
6.1 References.....	11
6.2 Author’s Address.....	12

0.5 Terminology and Document Conventions

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY" and "OPTIONAL" in this document are to be interpreted as described in 6.1.

Source code is shown in this typeface.

0.6 Revision History

The last named individual in this history is currently responsible for this document.

Revision	Date	Comments
Initial	January 2010	David Bosschaert (david.bosschaert@progress.com) Separated Aysnc Services from the general asynchronous communications RFP 124.
0.1	March 11 th , 2010	David Bosschaert (dbosscha@progress.com) Rewrote requirements section following focus group conference call (March 9 th)
0.2	September 21 st , 2013	Reactivate RFP focussing on asynchronous P2P execution, rather than PubSub and messaging, which is covered under other RFPs
0.3	September 26 th , 2013	Update Logo and Licence to match the current RFP template
0.4	October 24 th , 2013	Updates based on EEG review. Small tweaks to Application Domain. Made explicit the requirement for a generic mechanism to identify "async capable" services.
1.0	November 15 th , 2013	Discussed at virtual Face to Face. Revert requirement number changes and add the new requirement from version 0.4 as AS10. Accept all other changes and mark document as final and ready for vote.

1 Introduction

Over the last decade there has been a significant shift in computing hardware. Most systems now have access to multiple CPU cores, and can therefore only achieve optimal throughput when processing work on multiple threads. When mixed workloads are being processed there are further benefits to having multiple threads, computational jobs can continue while other jobs are blocked, for example waiting on IO devices.

Historically these improvements in hardware and throughput have been achieved through “Request level parallelism”. In this case multiple client requests (e.g. HTTP GET requests) are processed simultaneously on separate threads, but each request is executed in a single-threaded way. This works well – up to a point. If the individual requests are long-running, or result in multiple further requests, then these requests may proceed slowly, despite CPU resource being available.

To best leverage the capabilities of modern systems programmers are making more use of asynchronous execution. Some single threaded systems, such as JavaScript engines, make heavy use of asynchronous processing to allow multiple functions to appear to run concurrently. Several different mechanisms exist for exploiting asynchronous execution.

- Fire and forget – In the case that the return value of a method is not needed in subsequent execution, calling the method has no local side-effects, and need not complete before the caller returns, then it can be started and left running. Good tasks for this sort of execution model include
 - Submitting an “order” to a back-end system, which will be batch processed later
 - Logging to a central log store
- Blocking Promises – In some cases the return value from one or more methods is needed before execution can complete, however if these methods do not have a “happens before” dependency relationship then they can be executed asynchronously and potentially in parallel. Once submitted the caller receives a “promise” which can be queried to find out the eventual result. These are known as “blocking” promises because the initiating thread blocks until the result is available. Java’s Future/ExecutorService API is an example of this sort of asynchronous model. The Blocking Promises model has two significant advantages over the fire and forget model; it supports returning values and it allows the caller to wait until the asynchronous action has completed. Blocking the calling thread can, however, lead to the same sorts of problems seen with request-level parallelism. The calling thread is “lost” to the system until the necessary asynchronous tasks complete. If the asynchronous tasks also make use of blocking promises then the asynchronous threadpool can be exhausted, leading to deadlock.
- Promises with callbacks – JavaScript makes heavy use of promises, but due to JavaScript’s single-threaded execution model it is not possible for callers to block without deadlocking the whole process. As a result JavaScript promises work differently, rather than providing a blocking get method javascript promises are passed a function to execute on successful completion, or on failure. This callback mechanism ensures that deadlock does not occur, but still allows access to the return value of the asynchronous execution. The most popular promise implementations also provide ways to chain and aggregate promise returns, delaying callbacks until the relevant tasks are complete. Because of the way that promises with callbacks work they have typically been less popular in Java, where declaring an inner class as a callback is quite verbose. This is likely to change with the introduction of lambda expressions in Java 8.

In OSGi services are used as a way for loosely coupled bundles to interact, and represent a natural boundary for execution. Services declare an interface contract (which may be a Java class or interface), which can be referenced and called without having to directly load or instantiate the implementation. Many service calls are short, but some may be long-running, or delegate to remote implementations. In general OSGi services execute synchronously, and only offer asynchronous behaviour when it is explicitly built in to both the API and implementation. This obviously limits the use of asynchronous programming paradigms in OSGi.

The current Remote Services specification defines a set of properties to add to an OSGi service to enable it for remote communication across address spaces using existing distributed computing software systems, using synchronous request/response semantics. In this way Remote Services behave much like local OSGi services.

The synchronous semantics of Remote Services meet many requirements for distributed computing, but due to their higher latency and cost they can benefit even more from the capabilities of asynchronous communication. One need only look at frameworks such as AKKA to see how asynchronous protocols can be used to improve the simplicity and scaling performance of distributed systems. Given that Remote Services can benefit so greatly from asynchronous execution additional requirements should be included to ensure that asynchronous communication models and protocols can be leveraged where possible.

Requirements for an asynchronous programming model should be considered additional to the existing Remote Services Specification's request/response model, not as a replacement.

2 Application Domain

This section explores various aspects of adding support for asynchronous execution. Asynchronous execution typically is achieved via the introduction of a queuing mechanism for “tasks” which are pulled in and executed by one or more Threads. In the case of remote invocations, the task queue is often on the remote machine, allowing the request to be sent and executed without occupying a local Thread. These mechanisms are often also used to handle events, for example the OSGi Event Admin Service provides an asynchronous communication model.

Synchronous invocations are typically easier to program, but once a client makes a request, either local or remote, the client is blocked waiting for execution to complete and return control to the client. While asynchronous execution may be more complex to program, it offers many benefits and advantages.

For example, synchronous remote invocations depend on the availability of the network during request execution. If a client or server fails during the execution of a request, the request typically has to be resubmitted. This may not be a problem for some applications, where it's easy to re-create the request input. But for other applications, such as an ATM, gas pump, or electronic funds transfer, it may not be easy to recapture the input data and create another request message, and asynchronous protocols meet the requirement better. Even when it is possible to recreate a request message, it is not always easy to know at which point the server failed – i.e. whether or not an update was performed as a result of executing the request, and if so, whether performing the update a second time might cause data inconsistency. And in this case asynchronous protocols can also offer some advantages.

Synchronous invocations operate on a first-come, first-served scheduling mechanism (i.e. the computer has to process requests as they are made by the caller). This means that it's not easy to treat some invocations with higher priority than others, although this is a common application requirement (for example, a bank wants to process the outstanding \$1M deposits ahead of the \$10 deposits near the end of the banking day). As they have work queues, asynchronous processing engines can process work in an arbitrary order if they choose.

2.1 Asynchronous programming models

A variety of asynchronous programming models and frameworks are successfully used in enterprise applications today, including `ExecutorServices`, `Async EJBs`, `Async Servlets`, `Node.js`, `store-and-forward`, `pub-sub`, and `broadcast/multicast` to name a few. These programming models assume that a task is visible to a program using one or more asynchronous submission mechanisms (for example, `JMS`) and that the program is responsible for explicitly creating or retrieving a response using the API and then may act upon it in a way that is visible to another program using the same API.

For example, a store and forward system has one program submitting a message to a queue using a `SEND` or `SUBMIT` command, and another program retrieving the message from a queue using a `RECEIVE` or `DEQUEUE` command from the asynchronous programming model API. The sending program is responsible for packing, or serializing, the message, and the receiving program is responsible for unpacking, or de-serializing the message. (Some APIs define a wire format while others do not.)

In each case, management utilities are required to configure the capabilities of the asynchronous implementation being used so that they are able to reject work when overloaded, make best use of the resources available, and to identify, report and resolve any errors that may occur.

2.2 Mixture of programming models

Many enterprise applications require both synchronous and asynchronous execution models for different types of IT functions. For a reserved ticket purchase, for example, it may be necessary to synchronize the database update with the reply to the user to indicate the ticket was purchased, since only one person can have a given seat. For a book purchase, however, it may be sufficient to reply to the user that the order was received, and that it would be fulfilled later. Some of the fulfillment operations for a book order might also use synchronous communications, for example to debit inventory while packing the order for shipment.

2.3 Terminology + Abbreviations

3 Problem Description

The current OSGi programming model for communications among components and bundles is based on the OSGi service interface, which implies a synchronous semantic (i.e. the client invokes on the interface and waits for the reply), and language objects as parameters. These characteristics are typical of local invocations and distributed RPC and meet many requirements, but we want to extend these capabilities to support asynchronous invocation.

3.1 Asynchronous Services

We propose that the EEG evaluate options for specifying Asynchronous invocation of services – specifically the ability for a client to issue an invocation on a service interface without waiting for completion, and relying on a later notification or polling to check completion and retrieve results. For illustration, a low-level equivalent of such a framework is provided in J2SE by the Future interface. Other technologies (such as CORBA) provide asynchronous 'one-way' support on their remote interfaces. There are significant design considerations involved in selecting whether this may be defined within the “OSGi Services” architecture, and/or “Blueprint”, and/or Remote Services; and how a particular choice of solution relates to all three architectures.

4 Use Cases

The use cases describe situations in which requests are received and processed at some later time, which are typically the kind of use cases for which asynchronous execution is a better solution. Similarly, the use cases include situations in which the sending and receiving programs communicate indirectly using some type of intermediate store.

1. A web store receives an order for a book and submits the order immediately into a queue, letting the customer know the order was received. Later, perhaps hours or even days later, the order fulfillment program retrieves the order from the queue and processes it.
2. A bank receives an electronic funds transfer request and immediately puts the information into a file, and notifies the sender that the transfer was received. Later, probably seconds or minutes later, the banking system deposits (or withdraws) the transfer amount into (or from) the appropriate account.
3. A telecommunications company's network switch detects an error and publishes an event to report it. Later, probably seconds later, the network management system receives the event and displays a message on the console to alert the operator.
4. An electronics company places an order for parts to build PCs, potentially from different suppliers for CPUs, disks, displays, memory chips, etc. Each supplier acknowledges receipt of the order and notifies the electronics company of the likely shipment date and cost. Later, each supplier processes and ships the order, and notifies the electronics company the order has shipped. Even later, the electronics company receives and confirms all the shipments, ending the transaction with multiple parties, each of whom receives the confirmation separately.
5. A photography website accepts uploads from users, then performs facial recognition to automatically “tag” known users in the photos. When uploads complete they need to trigger the facial recognition, but should not block the sites ability to accept new upload requests.
6. In using an ATM or other point of sale device, transaction requests need to be captured whether the back office server is online or not. Transactions also need to be processed by the server once and only once.

5 Requirements

AS01 – The solution **MUST** provide a standard client-side API for making asynchronous invocations on existing, synchronous, OSGi services, where the invocation returns quickly and a return value can be obtained later.

AS02 – The solution **MUST** allow transparent delegation to services that are already implemented in an asynchronous fashion, therefore servicing the asynchronous requests through their own implementations.

AS03 – The solution **MAY** provide a synchronous client-side API to services which are implemented in an asynchronous fashion.

AS04 – The solution **MUST** allow for one-way (fire and forget) asynchronous services.

AS05 – The solution **MUST** support Promises, where invocations can be made that later return a value

AS06 – The solution **SHOULD** support callbacks when asynchronous executions complete, both successfully and unsuccessfully

AS07 – The solution **MUST** be applicable to both local OSGi Services as well as Remote OSGi Services.

AS08 – The solution **MUST** be fully backwards compatible with existing OSGi Service and Service Registry usage.

AS09 – The solution **SHOULD** be sympathetic to Java 8's lambda support, meaning callbacks should follow the Single Abstract Method principle where possible.

AS10 – The solution **MUST** define a mechanism that allows service providers to advertise an asynchronous mode of operation if they support it.

6 Document Support

6.1 References

- [1]. Bradner, S., Key words for use in RFCs to Indicate Requirement Levels, RFC2119, March 1997.
- [2]. Software Requirements & Specifications. Michael Jackson. ISBN 0-201-87712-0

6.2 Author's Address

Name	Eric Newcomer
Company	Progress Software
Address	14 Oak Park Drive, Bedford MA USA
Voice	+1 781 280 4000
e-mail	erickn@progress.com

Name	Giovanni Boschi
Company	Progress Software
Address	14 Oak Park Drive, Bedford MA USA
Voice	+1 781 280 4000
e-mail	giovanni.boschi@progress.com

Name	David Bosschaert
Company	Progress Software
Address	158 Shelbourne Road BallsBridge Dublin 4 Ireland
Voice	+353 1 637 2000
e-mail	dbosscha@progress.com

Name	Tim Ward
Company	Paremus Ltd
Address	
Voice	
e-mail	Tim.ward@paremus.com

6.3 End of Document