



RFP 158 - Distributed Eventing

Draft

15 Pages

Abstract

The OSGi specifications have described how to distribute OSGi services across VM boundaries now for a number of years. However OSGi services are synchronous in their nature. Many of today's business applications require asynchronous distributed communication mechanisms. While the OSGi Event Admin specification describes an asynchronous eventing model inside the Java VM this does not address event distribution to other VMs. This RFP aims to address the issue of Distributed Events in an OSGi context.

0 Document Information

0.1 License

DISTRIBUTION AND FEEDBACK LICENSE, Version 2.0

The OSGi Alliance hereby grants you a limited copyright license to copy and display this document (the "Distribution") in any medium without fee or royalty. This Distribution license is exclusively for the purpose of reviewing and providing feedback to the OSGi Alliance. You agree not to modify the Distribution in any way and further agree to not participate in any way in the making of derivative works thereof, other than as a necessary result of reviewing and providing feedback to the Distribution. You also agree to cause this notice, along with the accompanying consent, to be included on all copies (or portions thereof) of the Distribution. The OSGi Alliance also grants you a perpetual, non-exclusive, worldwide, fully paid-up, royalty free, limited license (without the right to sublicense) under any applicable copyrights, to create and/or distribute an implementation of the Distribution that: (i) fully implements the Distribution including all its required interfaces and functionality; (ii) does not modify, subset, superset or otherwise extend the OSGi Name Space, or include any public or protected packages, classes, Java interfaces, fields or methods within the OSGi Name Space other than those required and authorized by the Distribution. An implementation that does not satisfy limitations (i)-(ii) is not considered an implementation of the Distribution, does not receive the benefits of this license, and must not be described as an implementation of the Distribution. "OSGi Name Space" shall mean the public class or interface declarations whose names begin with "org.osgi" or any recognized successors or replacements thereof. The OSGi Alliance expressly reserves all rights not granted pursuant to these limited copyright licenses including termination of the license at will at any time.

EXCEPT FOR THE LIMITED COPYRIGHT LICENSES GRANTED ABOVE, THE OSGi ALLIANCE DOES NOT GRANT, EITHER EXPRESSLY OR IMPLIEDLY, A LICENSE TO ANY INTELLECTUAL PROPERTY IT, OR ANY THIRD PARTIES, OWN OR CONTROL. Title to the copyright in the Distribution will at all times remain with the OSGi Alliance. The example companies, organizations, products, domain names, e-mail addresses, logos, people, places, and events depicted therein are fictitious. No association with any real company, organization, product, domain name, email address, logo, person, place, or event is intended or should be inferred.

THE DISTRIBUTION IS PROVIDED "AS IS," AND THE OSGi ALLIANCE (INCLUDING ANY THIRD PARTIES THAT HAVE CONTRIBUTED TO THE DISTRIBUTION) MAKES NO REPRESENTATIONS OR WARRANTIES, EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, OR TITLE; THAT THE CONTENTS OF THE DISTRIBUTION ARE SUITABLE FOR ANY PURPOSE; NOR THAT THE IMPLEMENTATION OF SUCH CONTENTS WILL NOT INFRINGE ANY THIRD PARTY PATENTS, COPYRIGHTS, TRADEMARKS OR OTHER RIGHTS.

NEITHER THE OSGi ALLIANCE NOR ANY THIRD PARTY WILL BE LIABLE FOR ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF OR RELATING TO ANY USE OR DISTRIBUTION OF THE DISTRIBUTION.

Implementation of certain elements of this Distribution may be subject to third party intellectual property rights, including without limitation, patent rights (such a third party may or may not be a member of the OSGi Alliance). The OSGi Alliance is not responsible and shall not be held responsible in any manner for identifying or failing to identify any or all such third party intellectual property rights.

The Distribution is a draft. As a result, the final product may change substantially by the time of final publication, and you are cautioned against relying on the content of this Distribution. You are encouraged to update any implementation of the Distribution if and when such Distribution becomes a final specification.

The OSGi Alliance is willing to receive input, suggestions and other feedback ("Feedback") on the Distribution. By providing such Feedback to the OSGi Alliance, you grant to the OSGi Alliance and all its Members a non-exclusive, non-transferable,

worldwide, perpetual, irrevocable, royalty-free copyright license to copy, publish, license, modify, sublicense or otherwise distribute and exploit your Feedback for any purpose. Likewise, if incorporation of your Feedback would cause an implementation of the Distribution, including as it may be modified, amended, or published at any point in the future ("Future Specification"), to necessarily infringe a patent or patent application that you own or control, you hereby commit to grant to all implementers of such Distribution or Future Specification an irrevocable, worldwide, sublicenseable, royalty free license under such patent or patent application to make, have made, use, sell, offer for sale, import and export products or services that implement such Distribution or Future Specification. You warrant that (a) to the best of your knowledge you have the right to provide this Feedback, and if you are providing Feedback on behalf of a company, you have the rights to provide Feedback on behalf of your company; (b) the Feedback is not confidential to you and does not violate the copyright or trade secret interests of another; and (c) to the best of your knowledge, use of the Feedback would not cause an implementation of the Distribution or a Future Specification to necessarily infringe any third-party patent or patent application known to you. You also acknowledge that the OSGi Alliance is not required to incorporate your Feedback into any version of the Distribution or a Future Specification.

I HEREBY ACKNOWLEDGE AND AGREE TO THE TERMS AND CONDITIONS DELINEATED ABOVE.

0.2 Trademarks

OSGi™ is a trademark, registered trademark, or service mark of the OSGi Alliance in the US and other countries. Java is a trademark, registered trademark, or service mark of Oracle Corporation in the US and other countries. All other trademarks, registered trademarks, or service marks used in this document are the property of their respective owners and are hereby recognized.

0.3 Feedback

This document can be downloaded from the OSGi Alliance design repository at <https://github.com/osgi/design> The public can provide feedback about this document by opening a bug at <https://www.osgi.org/bugzilla/>.

0.4 Table of Contents

0 Document Information.....	2
0.1 License.....	2
0.2 Trademarks.....	3
0.3 Feedback.....	3
0.4 Table of Contents.....	3
0.5 Terminology and Document Conventions.....	4
0.6 Revision History.....	4
1 Introduction.....	5
2 Application Domain.....	5
2.1 Point-to-point/Queue semantics with current Event Admin Service.....	8
2.2 Existing approaches to distribute the Event Admin Service.....	8
2.3 Terminology + Abbreviations.....	8
3 Problem Description.....	9
4 Use Cases.....	11
4.1 Remote Service Administration.....	11
4.2 Config Administration.....	11
4.3 Distributed Event Based Interactions.....	12
4.4 Distributed Job Processing.....	12
5 Requirements.....	13

6 Document Support.....	14
6.1 References.....	14
6.2 Author's Address.....	14
6.3 End of Document.....	15

0.5 Terminology and Document Conventions

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY" and "OPTIONAL" in this document are to be interpreted as described in 6.1.

Source code is shown in this typeface.

0.6 Revision History

The last named individual in this history is currently responsible for this document.

Revision	Date	Comments
Initial	02/27/13	R Nicholson: Initial Cut
0.1	04.03.13	M. Schaaf; Added additional requirements and basic definitions
0.2	03/07/13	David Bosschaert, comments and small additions
0.3	03/03/13	R Nicholson - updated use-cases, requirements section (in response to BJ), added Abode use case & general re-work.
0.4	April, 2013	David Bosschaert, feedback from the F2F. Add use case from Carsten Ziegeler.
0.5	02.05.13	Marc Schaaf, Added another use case description for EDAs.
0.6	June, 2013	Process feedback from the Palo Alto F2F.

1 Introduction

This RFP presents the case for, and explores use cases and requirements of, a distributed event service.

As distributed events are increasingly seen as an essential component of a Cloud solution; this RFP compliments RFP 133 and RFC 183. This RFP also makes the case that RFC 185 (Data Transfer Objects) provide an excellent basis for cross-JVM messaging – be this Service property announcements or messages / events.

While suggesting that a distributed eventing service may enable a compelling foundation for distributed asynchronous programming models; the topic of asynchronous programming frameworks is out of scope of this RFP. While introduced, asynchronous RPC Services are also out of scope.

2 Application Domain

Distributed systems may be built using a number of different *interaction patterns*. Despite vocal proponents for each approach - it is increasingly clear that no one architectural solution is optimal in every context. Rather there is a continuous spectrum of interaction behaviors. If at all possible – these should ideally be supported in a consistent / coherent manner.

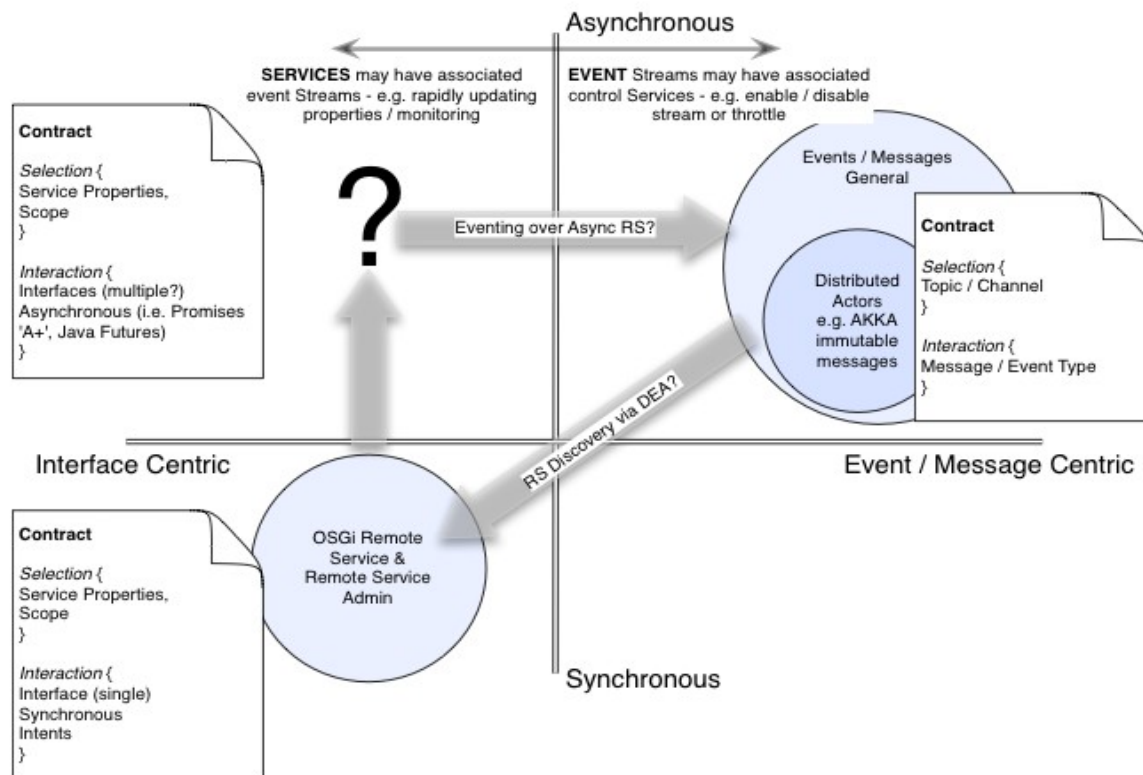


Figure 1: Types of distributed interaction

Synchronous RPC Services: The OSGi Alliance addressed the requirement for Remote Services via the *remote service* and *remote service admin specifications*: these specifications for synchronous RPC Services. In a dynamic environment (1..n) Services may be dynamically discovered, a sub-set of which (1..m where $m < n$) may be selected for use based on advertised Service properties. Service properties might be *Immutable* e.g. *Version*, *Location* or *Ownership* information – or *Mutable*: *reliability metrics*, *rating* or *cost metrics* – which may changing frequently in time.

It should be noted that:

- The RSA architecture is modular – allowing different Data Distribution and Discovery Providers to be used. This approach is extremely flexible. Some RSA implementations may choose to use a distributed P2P event substrate to provide Service discovery while other implementations use some form of static look-up service. Which ever is used - a coherent OSGi architecture is maintained. The use of an distributed Event Substrate for Service Discovery is one example of how RS/RSA and Distributed Eventing might interact.
- The current RSA specification does not address Service property updates: properties may change – and one does not necessarily wish to remove and re-add a Service because of this change. In more extreme cases, for volatile Service properties, one may wish to monitor these. Here a reference to the appropriate event streams might be advertised as Service properties. This scenario highlights a second potential relationship between RS/RSA and Distributed Eventing. (Note that it is planned to update the Remote Service Admin specification for Enterprise R6 to support Service property updates.)

Note the suggested RS/RSA enhancements are out of scope of this RFP – they are mentioned to illustrate the potential relationships between RS/RSA and a Distributed Eventing specification.

Asynchronous Services: Synchronous Services will block until a response is received from the local or remote service. However, on occasions it is preferable for the client to return from the call immediately and continue – some form of notification received at a future point in time – i.e. a 'promise' or a 'future'.

While a number of remote service data distribution providers (RSA) can – in principle support - asynchronous operation, there are currently no OSGi specifications to address support for asynchronous Services (local or remote). Such a specification is desirable as asynchronous Services are increasing popular in Cloud and Distributed environments – and increasingly the JavaScript development community (e.g. node.js). Work in the planned JavaScript RFC will look at implementations of asynchronous Service Registries.

As indicated in Figure 1 – in static environments - Asynchronous Services might be used as a mechanism to implement distributed events. This makes less sense in dynamic environments as some form of discovery mechanism is required – which is usually event based for scaling. So Distributed Eventing would more likely underpin RS/RSA.

This is an important area that requires OSGi Alliance specification work, work that clear relates to both Distributed Eventing & RS / RSA, but Asynchronous Services are out of scope of this current RFP; they are the topic of RFP 132.

Distributed Events / Messages:

Asynchronous Message / Event based approaches are increasingly the underpinnings of large scale distributed environments including 'Cloud'. In these distributed systems *Publishers* endpoints are decoupled from *Subscribers* endpoints; association is achieved via the *Topic* the *Publishers* and *Subscribers* choose to subscribe to a named Topic - and /or – a specific Message type.

Implementations vary considerably – and range from 'classic' enterprise message solutions - (e.g. JMS/AMQP) with centralized message brokers – to peer-to-peer de-centralized P2P solutions – e.g. 0MQ and the Object Management Group's (OMG) Data Distribution Service -
see http://www.omg.org/technology/documents/dds_spec_catalog.html

In principle asynchronous message based system provide the potential for greater scalability. However one cannot naively claim that asynchronous messaging will always scale more effectively than synchronous services: the performance characteristics are implementation dependent. An asynchronous messaging Service implemented via a central broker may introduce significant latency, throughput bottlenecks and points of architectural fragility. Whereas a dynamic Services approach with effective Service endpoint load balancing capabilities – would avoid these issues. However, correctly implemented P2P asynchronous message based systems will out perform both - with lower latency, higher throughput and increased resilience.

Due to the increased level of end-point de-coupling and potentially the use of parallel asynchronous pipelines, interaction contracts within message / event / based systems are more challenging. Unlike a Service centric approach - failure of Subscribers is not obvious to Publishers (or visa-versa).

Dependent upon the Capabilities of the Distributed Eventing provider – events may / may not / be durable – and in-order delivery of message / events may / may not / be possible.

- Broker based messaging solutions (i.e. JMS Brokers) typical rely on ACID transactions between publishers and the message broker, then the message broker and subscribers. Such solutions are typically used for chaining interactions between coarse grained services running on large compute servers – i.e. Mainframes / large Unix Systems in traditional Enterprise environments. However centralized brokers / ACID transactions represent bottlenecks and points of fragility: failing to efficiently scale in large distributed highly parallel asynchronous environments.
- Increasingly highly distributed / parallelized environments typical of 'Cloud' are using P2P messaging solutions with compensational / recovery / eventual consistency / based approaches to recovery. In such environments the components with a distributed system need to be idempotent as messages / events / may be re-injected in response to some form of timeout or failure notifications. In such environments aggregation points are still required to coordination at input (fan-out) and output (fan-in) boundaries of the parallel flows.

From a developer perspective 'Actors' are an increasingly popular asynchronous programming style. While driven by the the Scala / Akka community – Java Actor frameworks also exist – i.e. the kilim actor framework (<http://www.mahar.net/sriram/kilim/>). In these environments local asynchronous events (locally using a message / mailbox pattern) may be distributed between remote 'Actors' via a plug-able messaging layer; e.g. for Akka OMQ or via Camel / Camel plug-ins. An OSGi Distributed Eventing specification would provide a natural remoting substrate for 'Actor' / OSGi based applications.

2.1 Point-to-point/Queue semantics with current Event Admin Service

Some projects use the OSGi Service Registry Hooks to add point-to-point and/or queue messaging semantics to existing Event Admin Service implementations. This approach is working well for these projects and does not actually require a change to the Event Admin Service specification as it uses the hooks to only show the listeners that should receive the message to the Event Admin Service. While not distributed across remote frameworks such as design could also be relevant in a distributed context.

2.2 Existing approaches to distribute the Event Admin Service

A number of projects have successfully implemented a distribution-enabled Event Admin Service employing the existing OSGi API of the Event Admin Service to send events to remote clients. A master thesis was also written on the topic in 2009 by Marc Schaaf [4].

While this approach is very useful in certain situations, it has limitations which make the current Event Admin Service not generally applicable as a service for distributing events.

2.3 Terminology + Abbreviations

Messages: Are data structures. The structure could be as simple as the string 'BUY' or 'SELL' or it could be a complex structure like an object graph serialized to JSON. A message may be processed in isolation of other messages and have no consequence to the behavior of the node processing the message.

Events: Events are messages with additional 'significance' and 'relativity' - Reference - ['The Power of Events'](#)

'Significance' → the message should contain data describing the event it signifies & may prompt an action or change of behavior.

'Relativity' → time, causality and aggregation.

- Time is easy – this event happened on 3/4/2013.
- Causality – What caused this message?
- Aggregation – if event A signifies the activity of the set of e is an aggregation of the events B1, B2, B3... then A is an aggregation of the events B(i) – where $i = 1, 2, 3, \dots$

The term '*message*' and '*messaging*' will be used to include simple messages and events.

Common definitions for messaging systems include:

Queue: A messaging channel which delivers each message only to one receiver even if multiple receivers are registered, the message will only be delivered to one of them. If a receiver disconnects than all following messages are distributed between the remaining recipients. (It should be configurable that if no recipient is registered when a message is about to be delivered if the message is kept until a receiver is registered or if the message will be lost)

Topic: A publish and subscribe channel that delivers a message to all currently subscribed receivers. Therefore one message is broadcasted to multiple recipients. If no subscription of a receiver is available when a message is about to be delivered, the message is discarded. If a messaging client is disconnected for a period of time, it will miss all messages transferred during this period.

3 Problem Description

The OSGi Alliance has an elegant approach to services & remote-services model via which local services, perhaps expressed by DS & Blueprint, may be simply made visible to clients in remote OSGi frameworks.

However, unlike Remote Services, the OSGi Alliance has no coherent approach to the support of distributed messaging / events. Given the increased mindshare in development communities – this driven by Cloud Computing and use of Actor type patterns, this is an important omission and is the focus of this RFP.

This is doubly surprising given that a number of existing OSGi specifications would benefit from such specifications:

- RSA might leverage a Distributed Eventing implementations for Service Discovery Events – i.e. to Announce/Publish local Service endpoints and Subscribe/Discover Remote Services endpoint events.
- A version of local Event Admin might leverage Distributed Eventing to distribute Events to remote frameworks.
- ConfigAdmin – local ConfigAdmin services might be update by remote Configuration Events.

Hybrid Interactions & Contracts

As previously suggested synchronous Services and asynchronous events represent two ends of a continuous spectrum of interaction behaviors.

Examples include:

- Service announcements and discovery (i.e. Service Events) are already used by RSA – though the implementation is viewed as RSA specific and effectively isolated from the rest of the OSGi framework – unless one treats an entity as a Service. (e.g. example in Cloud EcoSystems – RFC-183).
- We may wish to associate an event channel with an advertised Service – for communicating rapidly changing properties.
- Alternatively we may wish to have a synchronous Service interaction with users of a Topic – perhaps to change the rate/throttle message flow or change recovery behavior.

Defining a generic approach to Contracts or Service Level Agreements is still mostly an area of research – especially for non-functional properties (NFP) – (see SLang, CQML etc).

However the start of a coherent strategy for OSGi is suggested by P Kriens & BJ Hargrave who have argued that – from a Service perspective – interactions should be expressed in-terms of *contracts*; each participant having its own role with respect to the interaction contract: i.e.

- what the participant is expected to provide
- what the participant can expect from the other participants.

P. Kriens has said: “A contract is just the agreed set of interactions between modules. With Service based interactions one tends to think in terms of interfaces... e.g. the CreditCheck service provides method calculateRating(). This is a simplistic contract between one provider (i.e. the credit rating provider) and the consumers; and it doesn't appear to support asynchronous interaction because invocation always originates from a consumer.

Instead consider a contract based on a group of interfaces; this is somewhat more powerful as each participant can provide some interfaces and consume other interfaces. For example a stock exchange: the exchange itself provides the OrderEntry, and other participants provide ExecutionListeners or MarketDataProviders. Hence, the 'contract' is not with a single Java interface but with a coherent collection of interfaces: in other words a package.

The 'contract' concept - expressed as a data transfer object (DTO) - may span JVM boundaries and provide a consistent approach for; synchronous (simple interface – DTO defines rich set of service properties), asynchronous (DTO defines multiple interfaces) and event based (DTO defines event / message format) based interactions.”

While not defining a 'Contract' – DTO's provide generic foundations upon which 'Contract' descriptions may be created.

Distributed Eventing should naturally comply with this philosophy, as for Distributed Eventing the interaction '*contract*' is simply a combination of:

- The structure of the message / event
- The Topic
- The SLA provided by the Distributed Eventing implementation

Data transfer object specification (RFC 185) provides a natural natural representation for the structural payload of a distributed message/event. Meanwhile R5 *Capabilities* provide the natural mechanism via which a user may selected the appropriate Distributed Eventing implementation with respect to required SLA: (in-order delivery, durable or transient etc).

Issues with current Event Admin

It should be noted that the following issues exist with the current Event Admin specification. There is no concept of 'contract' and the messages are untyped, so each participant has to continually work out what kind of message it has received, validate it, handle errors and missing info, work out what it should send in response.

- Current Event Admin only specifies how to send and receive events
- What to do after receiving an Event is unspecified...
- Current Event Admin events are maps, where the values can be anything - Java's instanceof operator to find out the type. Does this / should this / be modernized to be DTO centric?

This is fine if we don't want to go to the trouble of defining a contract for a particular interaction, but the risk is that modules become *more* tightly coupled because of hidden assumptions about the form of events they exchange. Also Event Admin is missing features such as the ability to send a point-to-point reply to a specific message, perhaps to a specific endpoint or subset of endpoints (perhaps via correlation IDs).

For these reasons it may not be possible to repurpose the existing Event Admin since it is already designed for a certain set of local use-cases, and there may be backwards compatibility concerns. Hence a completely new distributed eventing design may be required that might optionally replace or complement the local Event Admin service.

4 Use Cases

The following use cases are envisioned.

4.1 Remote Service Administration

The remote service administration specification specifies a modular approach to remote services via its plug-able data provider, discovery provider and topology manager design.

- It should be possible for a RSA discovery provider to be a simple local client of a generic distributed event service – subscribing to 'Service Events' of the required type.
- It should be possible for a Remoted Service properties to be mapped to a Distributed Eventing Topic to facilitate updates of these.

4.2 Config Administration

'Configuration' is usually considered to be a local concern. By this we mean the setting of local properties on entities within the local environment - via (in the case of OSGi) the Config Admin service. However distributed

configuration is required when a remote management entity needs to set / or update / a PID across a number of remote endpoints as in Cloud environments.

Efficiency, resilience & scalability considerations suggest that such interactions should be asynchronous in nature. Also if one desires to apply configurations to a group / population, rather than individual endpoints, then a publish / subscribe event based model is suggested.

Hence a distributed event service is required. A configuration client would subscribe to 'Configuration Events' - within the desired scope / topic - from the Distributed Event Service; passing received events on - into the local ConfigAdmin service. Meanwhile a configuration management entity may:

- Discovery available Configuration Topics
- Join to the appropriate Configuration Topic
- Publish a configuration change event to selected Topic.
- Consume and apply a configuration change event from selected Topic.

4.3 Distributed Event Based Interactions

Within a single OSGi framework - the event admin service may be used to inform components about actions (e.g. in content management system (CMS) - "page created", "page updated" actions.

As soon the application runs in a clustered environment a mechanism is required to distribute these events. As local events are used – it is natural to continue using an event centric approach in a the distribution environment. This allows co-location or distribution of components in the runtime environment without refactoring the application: as would be the case if for example the distributed version of the application used a messaging system.

By marking the OSGi events with an appropriate Scope or Topic (in the simple case - 'local' or 'distributed') propagation of events is controlled – i.e. whether they should be distributed in a cluster - or just locally as you don't want to distribute things like bundle events etc.

A similar application area is the realization of (distributed) event driven architectures. They typically process event data with various event processing agents which can be deployed on a cluster of machines to achieve the required processing capacity. The Agents typically communicate with the help of event channels or topics. This communication can be realized by the distributed eventing mechanism discussed here. It would allow an application to abstract from the technical details of the actual communication technology which would be hidden by the distributed eventing implementation. Furthermore the application deployment can become quite flexible as remote eventing could possibly be realized by the same eventing service that can also be used for local eventing. In comparison with the current Event Admin, functionality is missing to provide (a) some form of queue based communication as well as (b) the ability to specify the Quality of Service regarding communication reliability to ensure that e.g. events aren't lost due to a receiving components restart. Such functionality will be provided by the distributed eventing.

The Remote Services specification describes the use of intent labels for defining constraints on a distribution provider. These constraints can require the distribution provider to use a certain protocol or provide a certain Quality of Service (QoS) for the communication. This mechanism allows the QoS of communications to be managed without tying to a specific implementation or surfacing implementation detail complexity. It is likely that there are common QoSs between Remote Services and Distributed Eventing and therefore a consistent approach is highly desirable.

4.4 Distributed Job Processing

Applications often have the need for asynchronous processing of data, like creating a thumbnail, transcoding a video, replicating data from a staging to a production environment etc. These "jobs" vary in their resource usage, however they have in common to be executed exactly once. By giving each job type a topic for identification and a properties map for the payload, this job description is very similar to the current OSGi event object. However current event admin does not have the capability to deliver an event exactly once.

A distributed eventing could be used for such distributed job scenarios where configurations specify how a job (based on the topic) should be processed. There are different processing rules depending on the job type like this job can be processed in parallel, needs to be processed in the order of arrival, or only a single job should be processed at a time etc.

The distributed eventing would distribute the jobs based on such configuration and the capabilities of the available instances (RFC-183 could be used to propagate the capabilities) and also ensure that a job is processed by exactly one instance including failover handling etc.

The Apache Sling project has currently such an implementation based on OSGi events.

5 Requirements

DE010 – The solution **MUST** allow the sending of asynchronous messages to remote recipients.

DE012 – The solution **MUST** support a one-to-many, pub-sub/topic messaging semantic.

DE015 – The solution **MUST** support a queue messaging semantic.

DE020 – The solution **MUST** be independent of messaging technology used. This may be message broker based, peer-to-peer using a centralized approach or otherwise.

DE030 – The solution **MUST** allow implementations to advertise their supported Qualities of Service.

DE040 – The solution **MUST** provide a mechanism to select an Event Service provider based on its provided QoS.

DE042 – The solution **SHOULD** define a list of well-known QoS. Implementations **MUST NOT** be required to support all of these well known QoS.

DE045 – An implementation **MUST** be allowed to provide additional proprietary Qualities of Service.

DE047 – The solution **MUST** enable the message sender to specify the actual QoS used for sending a certain message.

DE050 – Events / Messages **MUST** be language agnostic – enabling a remote non-Java party to participate; e.g. C/C++ OSGi based agents.

DE055 – The solution **MAY** define a standard message encoding, for example using DTO's, JSON and/or other technology if appropriate.

DE060 – The solution **MUST** provide the means for point-to-point based communications for example to allow replies to specific messages – an event targeted to a specific node.

DE080 – The solution **MUST** provide the means to obtain information on the sender of an event e.g. bundleID, Framework UUID, SubSystem name. This information **MAY** be incomplete if the message didn't originate in an OSGi framework.

DE085 – The solution **SHOULD** provide the means to discover available Topics and Queues..

DE087 – The solution **MUST** allow certain Topics and Queues to not be advertised in the discovery mechanism.

DE088 – The solution **SHOULD** allow certain messages to be hidden from potential receivers.

DE090 – The solution **SHOULD NOT** prevent an implementation from providing a basic distribution solution for the existing Event Admin. While this will not provide all features of a Distributed Eventing solution, it is shown to be useful in certain contexts.

6 Document Support

6.1 References

- [1]. Bradner, S., Key words for use in RFCs to Indicate Requirement Levels, RFC2119, March 1997.
- [2]. Software Requirements & Specifications. Michael Jackson. ISBN 0-201-87712-0.
- [3]. The Power of Events'. D. C. Luckham. Boston, MA, USA: Addison-Wesley Longman Publishing Co., Inc., 2001.
- [4]. Extending OSGi by Means of Asynchronous Messaging - Master Thesis, Marc Schaaf, September 2009, University of Applied Sciences and Arts Hannover.
http://schaaf.es/docs/master_thesis_marc_schaaf_Extending_OSGi_by_Means_of_Asynchronous_Messaging.pdf

6.2 Author's Address

Name	Richard Nicholson
Company	Paremus Ltd
Address	107-111 Fleet Street London
Voice	
e-mail	richard.nicholson@paremus.com

Name	Marc Schaaf
Company	
Address	
Voice	
e-mail	marc@marc-schaaf.de

Name	David Bosschaert
Company	Red Hat
Address	
Voice	
e-mail	david@redhat.com

Name	Carsten Ziegeler
Company	Adobe
Address	
Voice	
e-mail	cziegele@adobe.com

Name	Graham Charters
Company	IBM
Address	
Voice	
e-mail	

6.3 End of Document