



OSGiTM
Alliance

RFP 179 - Compute Management Service

Final

9 Pages

Abstract

A large variety of ways exist today to create, run and manage computing instances, whether this be in the cloud, on a computing grid or on a container-based environment. While physically highly diverse, managing these on a logical level can be seen as similar. This RFP seeks to define a common service to facilitate managing compute nodes across a wide variety of platforms.

0 Document Information

0.1 License

DISTRIBUTION AND FEEDBACK LICENSE, Version 2.0

The OSGi Alliance hereby grants you a limited copyright license to copy and display this document (the "Distribution") in any medium without fee or royalty. This Distribution license is exclusively for the purpose of reviewing and providing feedback to the OSGi Alliance. You agree not to modify the Distribution in any way and further agree to not participate in any way in the making of derivative works thereof, other than as a necessary result of reviewing and providing feedback to the Distribution. You also agree to cause this notice, along with the accompanying consent, to be included on all copies (or portions thereof) of the Distribution. The OSGi Alliance also grants you a perpetual, non-exclusive, worldwide, fully paid-up, royalty free, limited license (without the right to sublicense) under any applicable copyrights, to create and/or distribute an implementation of the Distribution that: (i) fully implements the Distribution including all its required interfaces and functionality; (ii) does not modify, subset, superset or otherwise extend the OSGi Name Space, or include any public or protected packages, classes, Java interfaces, fields or methods within the OSGi Name Space other than those required and authorized by the Distribution. An implementation that does not satisfy limitations (i)-(ii) is not considered an implementation of the Distribution, does not receive the benefits of this license, and must not be described as an implementation of the Distribution. "OSGi Name Space" shall mean the public class or interface declarations whose names begin with "org.osgi" or any recognized successors or replacements thereof. The OSGi Alliance expressly reserves all rights not granted pursuant to these limited copyright licenses including termination of the license at will at any time.

EXCEPT FOR THE LIMITED COPYRIGHT LICENSES GRANTED ABOVE, THE OSGi ALLIANCE DOES NOT GRANT, EITHER EXPRESSLY OR IMPLIEDLY, A LICENSE TO ANY INTELLECTUAL PROPERTY IT, OR ANY THIRD PARTIES, OWN OR CONTROL. Title to the copyright in the Distribution will at all times remain with the OSGi Alliance. The example companies, organizations, products, domain names, e-mail addresses, logos, people, places, and events depicted therein are fictitious. No association with any real company, organization, product, domain name, email address, logo, person, place, or event is intended or should be inferred.

THE DISTRIBUTION IS PROVIDED "AS IS," AND THE OSGi ALLIANCE (INCLUDING ANY THIRD PARTIES THAT HAVE CONTRIBUTED TO THE DISTRIBUTION) MAKES NO REPRESENTATIONS OR WARRANTIES, EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, OR TITLE; THAT THE CONTENTS OF THE DISTRIBUTION ARE SUITABLE FOR ANY PURPOSE; NOR THAT THE IMPLEMENTATION OF SUCH CONTENTS WILL NOT INFRINGE ANY THIRD PARTY PATENTS, COPYRIGHTS, TRADEMARKS OR OTHER RIGHTS.

NEITHER THE OSGi ALLIANCE NOR ANY THIRD PARTY WILL BE LIABLE FOR ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF OR RELATING TO ANY USE OR DISTRIBUTION OF THE DISTRIBUTION.

Implementation of certain elements of this Distribution may be subject to third party intellectual property rights, including without limitation, patent rights (such a third party may or may not be a member of the OSGi Alliance). The OSGi Alliance is not responsible and shall not be held responsible in any manner for identifying or failing to identify any or all such third party intellectual property rights.

The Distribution is a draft. As a result, the final product may change substantially by the time of final publication, and you are cautioned against relying on the content of this Distribution. You are encouraged to update any implementation of the Distribution if and when such Distribution becomes a final specification.

The OSGi Alliance is willing to receive input, suggestions and other feedback ("Feedback") on the Distribution. By providing such Feedback to the OSGi Alliance, you grant to the OSGi Alliance and all its Members a non-exclusive, non-transferable,

worldwide, perpetual, irrevocable, royalty-free copyright license to copy, publish, license, modify, sublicense or otherwise distribute and exploit your Feedback for any purpose. Likewise, if incorporation of your Feedback would cause an implementation of the Distribution, including as it may be modified, amended, or published at any point in the future ("Future Specification"), to necessarily infringe a patent or patent application that you own or control, you hereby commit to grant to all implementers of such Distribution or Future Specification an irrevocable, worldwide, sublicenseable, royalty free license under such patent or patent application to make, have made, use, sell, offer for sale, import and export products or services that implement such Distribution or Future Specification. You warrant that (a) to the best of your knowledge you have the right to provide this Feedback, and if you are providing Feedback on behalf of a company, you have the rights to provide Feedback on behalf of your company; (b) the Feedback is not confidential to you and does not violate the copyright or trade secret interests of another; and (c) to the best of your knowledge, use of the Feedback would not cause an implementation of the Distribution or a Future Specification to necessarily infringe any third-party patent or patent application known to you. You also acknowledge that the OSGi Alliance is not required to incorporate your Feedback into any version of the Distribution or a Future Specification.

I HEREBY ACKNOWLEDGE AND AGREE TO THE TERMS AND CONDITIONS DELINEATED ABOVE.

0.2 Trademarks

OSGi™ is a trademark, registered trademark, or service mark of the OSGi Alliance in the US and other countries. Java is a trademark, registered trademark, or service mark of Oracle Corporation in the US and other countries. All other trademarks, registered trademarks, or service marks used in this document are the property of their respective owners and are hereby recognized.

0.3 Feedback

This document can be downloaded from the OSGi Alliance design repository at <https://github.com/osgi/design> The public can provide feedback about this document by opening a bug at <https://www.osgi.org/bugzilla/>.

0.4 Table of Contents

0 Document Information.....	2
0.1 License.....	2
0.2 Trademarks.....	3
0.3 Feedback.....	3
0.4 Table of Contents.....	3
0.5 Terminology and Document Conventions.....	4
0.6 Revision History.....	4
1 Introduction.....	4
2 Application Domain.....	5
2.1 Terminology + Abbreviations.....	5
3 Problem Description.....	5
4 Use Cases.....	6
4.1 Docker cluster on local machine.....	6
4.2 Private cloud + Azure VM = hybrid.....	6
4.3 Multiple target platforms.....	6
4.4 Dynamic Cluster Scaling.....	7
4.5 Pets vs Cattle.....	7

5 Requirements.....	7
5.1 Persistent Volumes.....	8
5.2 Notifications.....	8
6 Document Support.....	8
6.1 References.....	8
6.2 Author's Address.....	9
6.3 End of Document.....	9

0.5 Terminology and Document Conventions

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY" and "OPTIONAL" in this document are to be interpreted as described in 6.1.

Source code is shown in this typeface.

0.6 Revision History

The last named individual in this history is currently responsible for this document.

Revision	Date	Comments
Initial	December, 2015	David Bosschaert < bosschae@adobe.com > initial version.
0.1	January, 2016	David Bosschaert, feedback from Madrid F2F
0.2	May, 2016	David Bosschaert, updates following discussion at the Chicago F2F
0.3	May, 2016	David Bosschaert, update to requirement NO0020 following feedback from Todor Boev.
Final	May, 2016	David Bosschaert, mark as final for vote.

1 Introduction

OSGi RFP 133 (http://osgi.org/bugzilla/show_bug.cgi?id=114) provides the foundation requirements of many cloud-related activities in the OSGi Alliance. It includes requirements around management of nodes in a cloud environment where a node runs an OSGi framework. This RFP seeks to elaborate on the cloud environment

management requirements, expanding them to more general container-based environments as well as adding support for nodes that can run any type of application, such as a database or load balancer.

2 Application Domain

The domain of this RFP is Cloud computing, Container-based computing and Grid computing. Effectively any environment where compute nodes or instances can be created programmatically. The RFP seeks to define requirements for a common API specification to manage Compute nodes across these environments.

Existing solutions in this area include, but are not limited to:

- Apache JClouds : <http://jclouds.apache.org>
- OpenStack: <http://www.openstack.org>
- Cloud Foundry: <https://www.cloudfoundry.org>
- TOSCA: <http://docs.oasis-open.org/tosca/>
- CAMP: https://www.oasis-open.org/committees/tc_home.php?wg_abbrev=camp
- Open Container Initiative: <https://www.opencontainers.org/>

2.1 Terminology + Abbreviations

Compute Node – an entity that can execute code. Either a physical computer, virtual machine or lightweight container. The entity provides an operating system, such as Microsoft Windows, Linux, other unix or similar. The node is connected to a IP-based network.

3 Problem Description

Many cloud vendors exist today that provide the capability to create compute nodes allowing the user to run applications on these vendor-provided environments. Compute nodes can generally be created using a remote API, web interface or command-line tool.

More recently container-based computing is becoming popular, docker is currently dominant in this space but other solutions such as rkt (Rocket) are also emerging. With container-based computing not an entire Virtual Machine is created, but rather a container to run a single program is created. While containers are isolated, the operating system and baseline functionality such as programming language support is often shared with other containers running on the same physical or virtual underlying platform.

While some standardization is happening on the virtual machine level, e.g. openstack, these standards don't apply to container based solutions. In the Java world some opensource projects exist that try to address this issue, such as Apache jclouds, however the biggest issue here is that as jclouds provides its own API it also needs to provide all of the bindings to the technologies. This means that jclouds supports an enormous amount of technologies, where not all of these are equally well maintained.

By defining an OSGi service API, implementors can decide to only support one target: the one they care about. Multiple targets can be supported at runtime by combining multiple API implementors, each of which will be represented, for example, by a service in the service registry. Having a specification-defined API supports this federated approach.

4 Use Cases

4.1 Docker cluster on local machine

Bob wants to run an application which is composed of a number of containers using Docker on his local machine. In the future he wants to deploy the same application to a docker-enabled cloud provider, but he does not know yet which one. Bob is looking to develop his application in such a way that he does not need to rewrite the container logic once he has decided what his ultimate target runtime is going to be.

4.2 Private cloud + Azure VM = hybrid

Rosanne needs to create an application that runs in a private cloud. However she needs to cater for scenarios where the capacity of the private cloud is not sufficient. In such cases she needs to be able to add public cloud Vms provided by Microsoft Azure to the application creating a hybrid deployment. She wants to keep the application logic simple and would like a single programming model that works for both the private cloud as well as in Azure.

4.3 Multiple target platforms

Harry needs to run an application that is comprised of a number of docker-based Vms in a cloud container platform such as Amazon ECS. However, for demonstration purposes he would also like to run this application on his local laptop using local docker containers. Harry wants to share his application logic across these two environments, he does not want to create different implementations of his work, one that runs on local docker and the other on Amazon ECS.

4.4 Dynamic Cluster Scaling

William runs his software in a cluster of machines. However, during certain times of the week, the cluster is much heavier loaded than others. William therefore installs a load checker on each machine and if the load goes over a certain limit then the software automatically creates a new node which is then added to the cluster. If the systems go below a certain limit the nodes are deleted again.

4.5 Pets vs Cattle

Zoe needs to run 4 Mongo DB nodes. She therefore creates 4 Elastic Block Storage (EBS) volumes and creates 4 compute nodes that attach these EBS volumes. She then starts the first mongo instance and initializes the support for replica sets. She uses the DNS name that is associated with the volume to mark the initial leader. She repeats this for the other 3 instances. The replica set is therefore made of the DNS names of the volumes.

She then marks the compute nodes to always have one instance running.

5 Requirements

CM0010 – The solution must provide a mechanism to create and destroy compute nodes.

CM0020 – It must be possible to implement the solution for existing compute platforms, including, but not limited to, cloud-based platforms, grid computing-based platforms and/or container-based platforms.

CM0030 – The solution must provide a mechanism to specify the root image to run on the compute node.

CM0040 – The solution must provide a mechanism to specify compute node parameters such as the amount of memory, cpu and exposed network ports.

CM0050 – The solution must provide a means to specify environment variables to be set in the compute node.

CM0060 – The solution must provide a mechanism to run an executable in the compute node once available.

CM0070 – The solution must support listing and querying of compute nodes.

CM0080 – The solution must support specifying that more than one node of a given definition be created.

CM0090 – The solution must allow the number of requested nodes of a given definition to be changed after the initial launch of the associated compute node(s).

CM0100 – The solution must support assigning logical names to nodes that can be used to address them.

CM0110 – The solution must allow integration with security features offered by the compute platform.

CM0120 – The solution must provide support for load balancers across a cluster of nodes.

CM0130 – The solution must provide the external address of a load balancer, if available.

CM0140 – The solution must support addressing a load balancer with a logical name.

CM0150 – The solution must provide runtime management information, for example via DTOs.

CM0160 – The solution must allow the effects of the service to follow a different lifecycle than the runtime of Compute Management Service.

CM0170 – The solution must be able to discover compute nodes created during previous runs of the Compute Management Service.

CM0180 – It must be possible to provide secrets to the computer node in a secure way

CM0190 – It must be possible to specify whether compute nodes have a private or public IP address or both.

CM0200 – It must be possible to create compute nodes which are not directly accessible by the creator, such as nodes that are connected to a different network than the creator.

5.1 Persistent Volumes

PE0010 – The solution must be able to create, list, and delete persistent volumes like Amazon Elastic Storage Blocks or Google's Persistent volumes

PE0020 – It must be possible to attach a persistent volume to a compute node when the node is created

PE0030 – It must be possible to have a unique DNS name for a compute node that is based on the attached persistent volume name.

5.2 Notifications

NO0010 – It must be possible to get notifications when nodes are created, are stopped, or die

NO0020 – It must be possible to receive notifications of the various states that a compute node goes through, including the following states: CREATED, WAITING, PENDING, STARTING, ACTIVE, STOPPING, STOPPED, FAILED. Backing platforms may not support all states. Implementations are not required to provide notifications for states not supported by the backing platform.

6 Document Support

6.1 References

- [1]. Bradner, S., Key words for use in RFCs to Indicate Requirement Levels, RFC2119, March 1997.

[2]. Software Requirements & Specifications. Michael Jackson. ISBN 0-201-87712-0

6.2 Author's Address

Name	David Bosschaert
Company	Adobe Systems
Address	
Voice	
e-mail	bosschae@adobe.com

Name	Peter Kriens
Company	aQute
Address	
Voice	
e-mail	

6.3 End of Document